


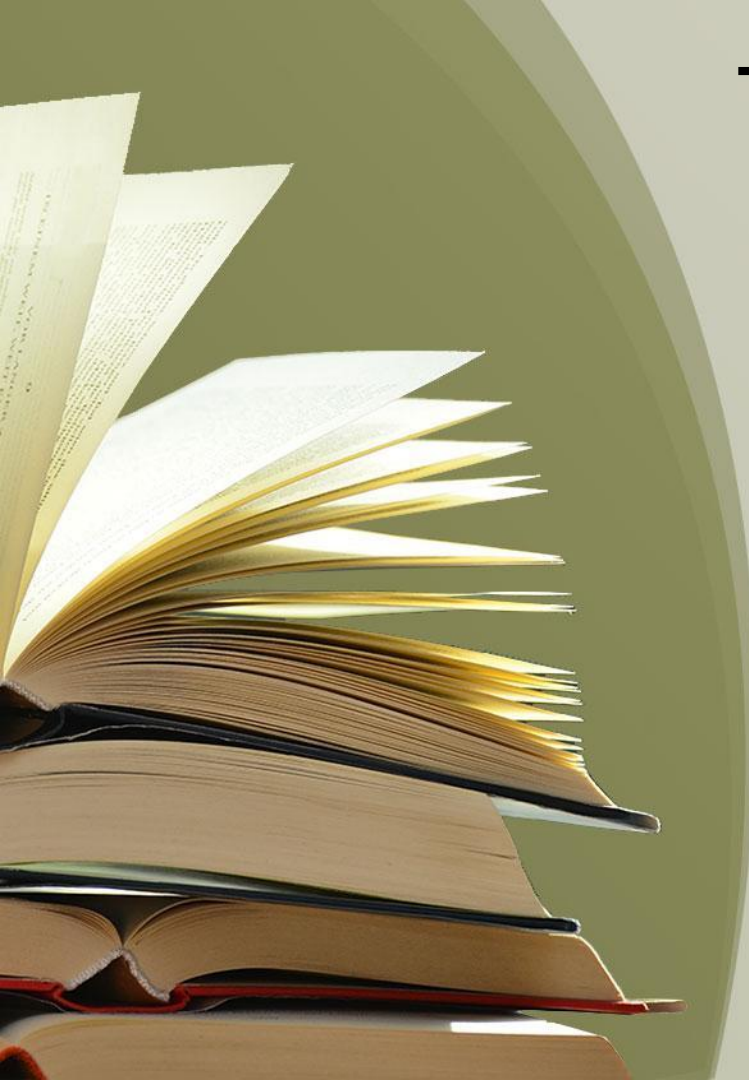
# Real Relevant Reviews (R<sup>3</sup>)

*Only The Most Helpful Product Reviews*

W266 Section 003

Jen Darrouzet & Bradley Nott





# The R<sup>3</sup> Project

1

## The Problem

Too many online reviews

2

## The Data

8 Million book reviews from Amazon.com

3

## The Approach

Helpful Vote Arrival Rates + Neural Nets

4

## The Results

To be revealed...


5

## The Next Steps

Real, relevant, crowd-sourced advice for buyers


# $R^3$ System






# The Problem

Need an efficient method to  
make product reviews useful



# The Data


8 Million book reviews from  
Amazon.com (2006-2014) that  
contain a helpful vote tally



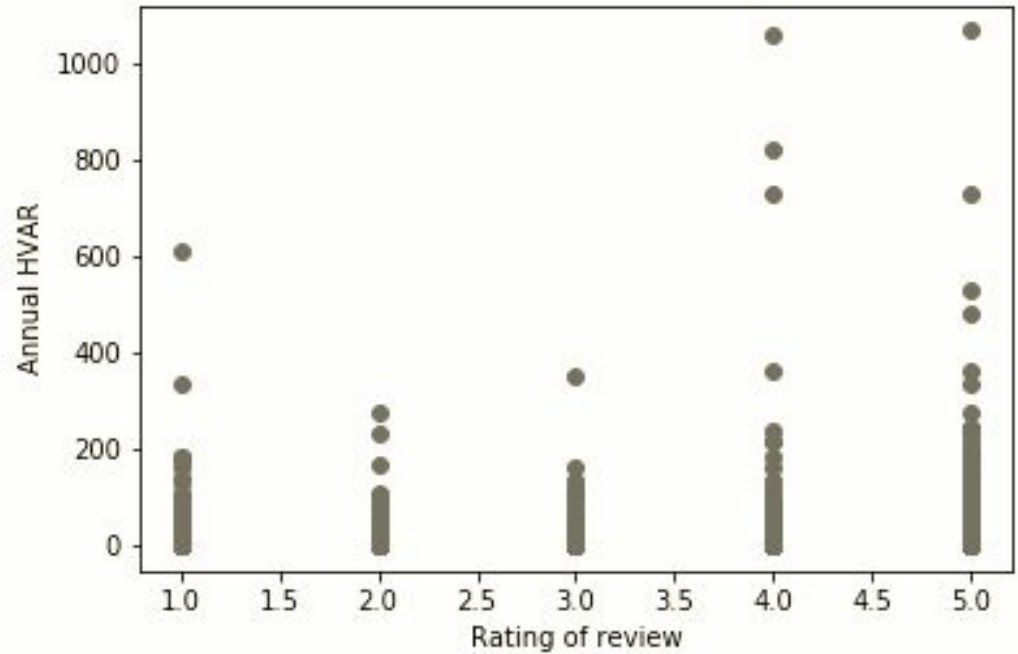
# The Approach

Transform the helpful vote tally  
into a metric that enables  
comparisons



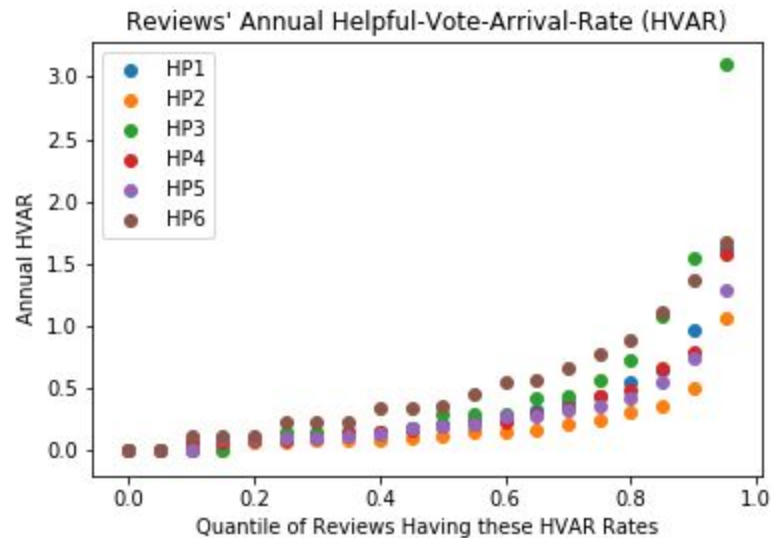



Critical Reviews Can Be Helpful, Too



# How to Train for Helpfulness

Counts? Proportions? Apples-to-apples?





# Experiments


- Data is noisy and the signal is rare
- What is “means” to be helpful varies
- Difficult to be confident about signal
- Smaller models have a hard time with variations in review length

# The Results



# Classifier Performance

Predicting helpfulness of a review, given review text and 1-5★ rating



## Modeling Notes

None of the models fit the noisy data well.  
Sorting by probability of helpfulness yields good top N.  
**Asking shoppers if we got it right could help, too.**

- |                              |   |
|------------------------------|---|
| <b>1 tf.idf &amp; LogReg</b> | Super fast; unlimited vocab was 1M+                       |
| <b>2 Bi-LSTM</b>             | Insanely slow (used alt CuDNNLSTM); 200 tokens            |
| <b>3 BERT</b>                | Medium speed; 200 tokens; 30K vocab; extra benefit of NSP |

## BERT's added benefit

If BERT is substantially same, **get relevance for "free"**



# Query Relevance

with BERT's Next Sentence Prediction (NSP)

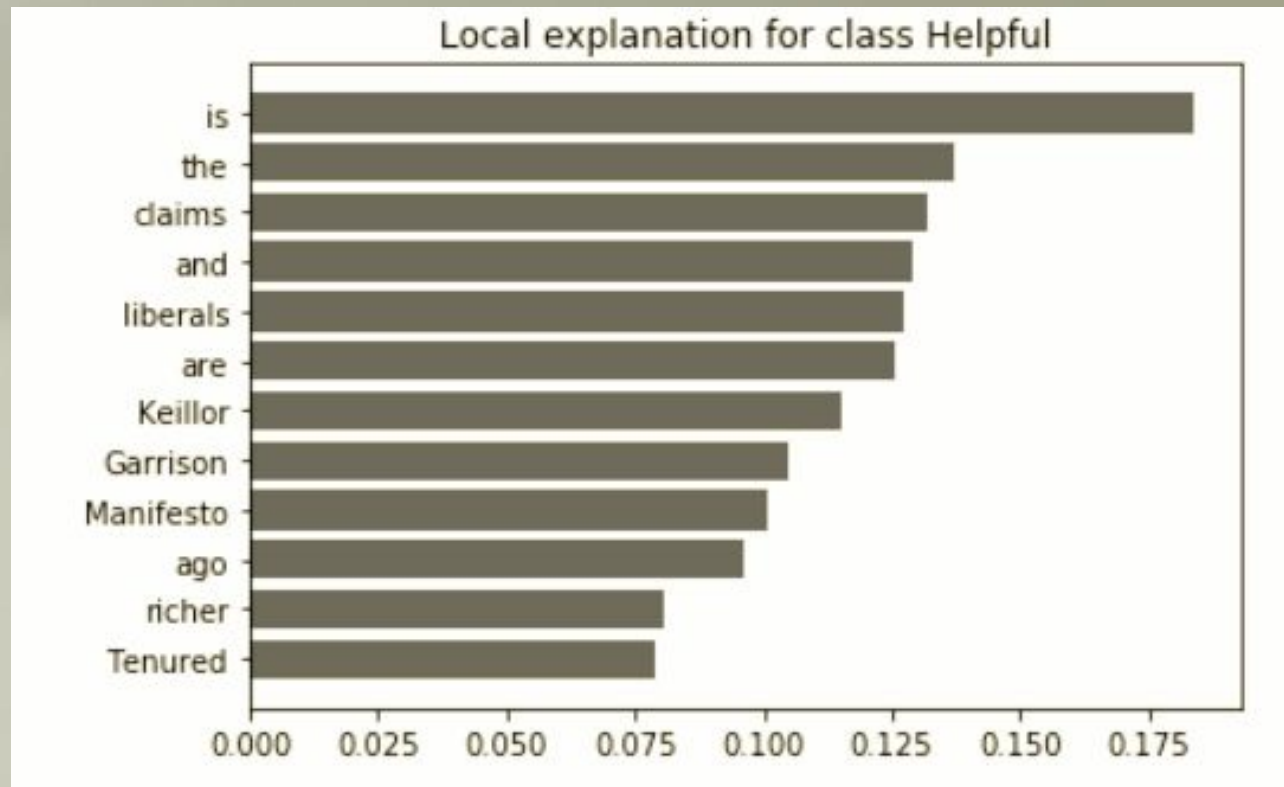
Query: *Is Water for Elephants appropriate for a teenager headed off to veterinary school?*

4★: ...Jacob's parents died just as he was about to write his final exams in **veterinary school** at Cornell University....**because of his vet training**, he ends up hired on to care for the horses and other animals...I read it quickly and found it very sweet and entertaining; **recommended for teens and adults**... Caveat: **there are descriptions of mistreatment and abuse of animals** in this book. If that is unpalatable to you, you should skip this book...

2★: Jacob, now an old man, begins the book by going back into his youth to the time **he was attending vet school**. ... He winds up with a circus and becomes the circus vet, although, **he rarely spent any time with or working on the animals**. **He just stood and watched animals be slaughtered and then felt bad** that he did not do anything...

# Helpfulness indicator

LIME helped identify that reviews with claims were >80% helpful





# Claims

Not just for nonfiction reviews, either

**BEST** Someone with an impressive array of murderous skills is systematically assassinating right-wing political figures one after another. The FBI and the CIA, who suspect that the killer comes from somewhere in their ranks, join forces to try to track down the assassin. An imprisoned former CIA traitor, who claims to know the identity of the assassin, wants to trade that information



# Next steps

Tune  $R^3$  for identified model weaknesses (e.g. length)

- ❑ Customer tolerance for review length may be product-category dependent
  - ❑ 📖📚 Book lovers embrace long-form text
  - ❑ 📌💄 Beauty product shoppers could prefer short text/user-submitted photos

Evaluate  $R^3$ 's ranked reviews using IR metrics like Mean Average Precision (MAP)

- ❑ Could not automatically measure precision/recall as dataset did not have relevance labels

Optimize  $R^3$  for shopper happiness (both **useful and fast**)

- ❑ Once we can reliably deliver a TOP N HELPFUL + RELEVANT review list, SPEED is paramount
  - ❑ Can we process new reviews immediately upon submission? (Important for new releases.)
  - ❑ What is the fastest BERT model to compare most helpful reviews to a shopper query?
  - ❑ Confident we can pre-calculate a BERT representation of each incoming review such that shoppers' query-time comparisons are fast...but how fast is the question.

